



CSE439-Introduction to Deep Learning

Salih Can AYDOĞDU

191805061

#Contains both the first and second parts of the project#

Detection of Diabetes Using Deep Learning

Diabetes is a metabolic disorder that occurs when the body is unable to produce a sufficient amount of insulin or effectively utilize it, leading to elevated levels of blood sugar. Early diagnosis and appropriate treatment can help control diabetes. In this context, the use of deep learning applications holds significant potential for effectively diagnosing diabetes and improving the treatment process. This research aims to develop a diabetes diagnostic model based on deep learning using clinical and physical data. To achieve this goal, various sampling techniques will be applied to the existing dataset, followed by the application of traditional machine learning algorithms and methods such as Long Short-Term Memory (LSTM), Gated Recurrent Units (GRU), and complex neural network architectures for diabetes diagnosis. The performance evaluation of this study will utilize the widely used Pima Indians Diabetes dataset. The obtained results will be used to assess whether this deep learning-based model is an effective tool for diabetes diagnosis.

1. The dataset used and its review

In this project I developed for the diagnosis of diabetes, I use the Pima Indians Diabetes Dataset, which is widely used around the world. If I examine the mentioned dataset in general;

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

Image 1.0

This dataset consists of a total of 768 observations, each representing an individual. Each observation includes information about a person's health characteristics and diabetes status. The features in the dataset are as follows:

1. **Pregnancies**: This feature represents the number of pregnancies a person has experienced. Its minimum value is 0 and its maximum value is 17.
2. **Glucose**: This feature expresses the glucose level in a person's blood in milligrams per deciliter (mg/dL).
3. **Blood Pressure**: This feature represents a person's blood pressure and is expressed in millimeters of mercury (mm Hg).
4. **Skin Thickness**: This feature represents the thickness of a person's skin in millimeters.
5. **Insulin**: This feature represents a person's insulin level and is expressed in microunits per milliliter (MicroU/mL).
6. **BMI (Body Mass Index)**: This feature represents a person's body mass index, calculated using the formula: $\text{Weight (kg)} / \text{Height (m)}^2$.
7. **Diabetes Pedigree Function**: This feature represents the value of a function based on the patient's diabetes history in the family tree.
8. **Age**: This feature represents a person's age. The youngest person is 21 years old and the oldest person is 81 years old. The overall average is 33.240885.
9. **Outcome**: This feature indicates whether a person has diabetes or not. 0 represents the absence of diabetes, and 1 represents the presence of diabetes. Approximately 34.90% of participants in the dataset are patients

2. Application of Dataset Balancing Methods and Their Interactions

In machine learning applications, success is generally proportional to the quality and balance of the dataset on which the model is trained. However, real-world datasets often tend to be imbalanced, meaning that one class may have significantly more examples than another. This imbalance can negatively impact the model's training process and lead to misleading results. Therefore, in machine learning applications, resorting to dataset balancing strategies has become inevitable.

When I examine the distribution of the Outcome feature in our dataset by using BMI and Glucose characteristics as independent variables, the distribution of the Outcome feature is as follows;




Image 2.0

As seen in the image 2.0, there is an imbalance in the data distribution in our outcome feature.

Feature	Before Sampling 0	Before Sampling 1	Sampling Method	After Sampling 0	After Sampling 1
Outcome	500	268	Before Sampling	500	268
Outcome	500	268	SMOTE	500	500
Outcome	500	268	KMeansSMOTE	500	502
Outcome	500	268	ADASYN	500	474
Outcome	500	268	SVM SMOTE	500	500
Outcome	500	268	RandomOverSampler	500	500
Outcome	500	268	BorderlineSMOTE	500	500
Outcome	500	268	EditedNearestNeighbours	240	268
Outcome	500	268	AIKNN	199	268
Outcome	500	268	InstanceHardnessThreshold	275	268
Outcome	500	268	NearMiss	268	268
Outcome	500	268	NeighbourhoodCleaningRule	250	268
Outcome	500	268	OneSidedSelection	422	268
Outcome	500	268	RandomUnderSampler	268	268
Outcome	500	268	TomekLinks	445	268

Image 2.1

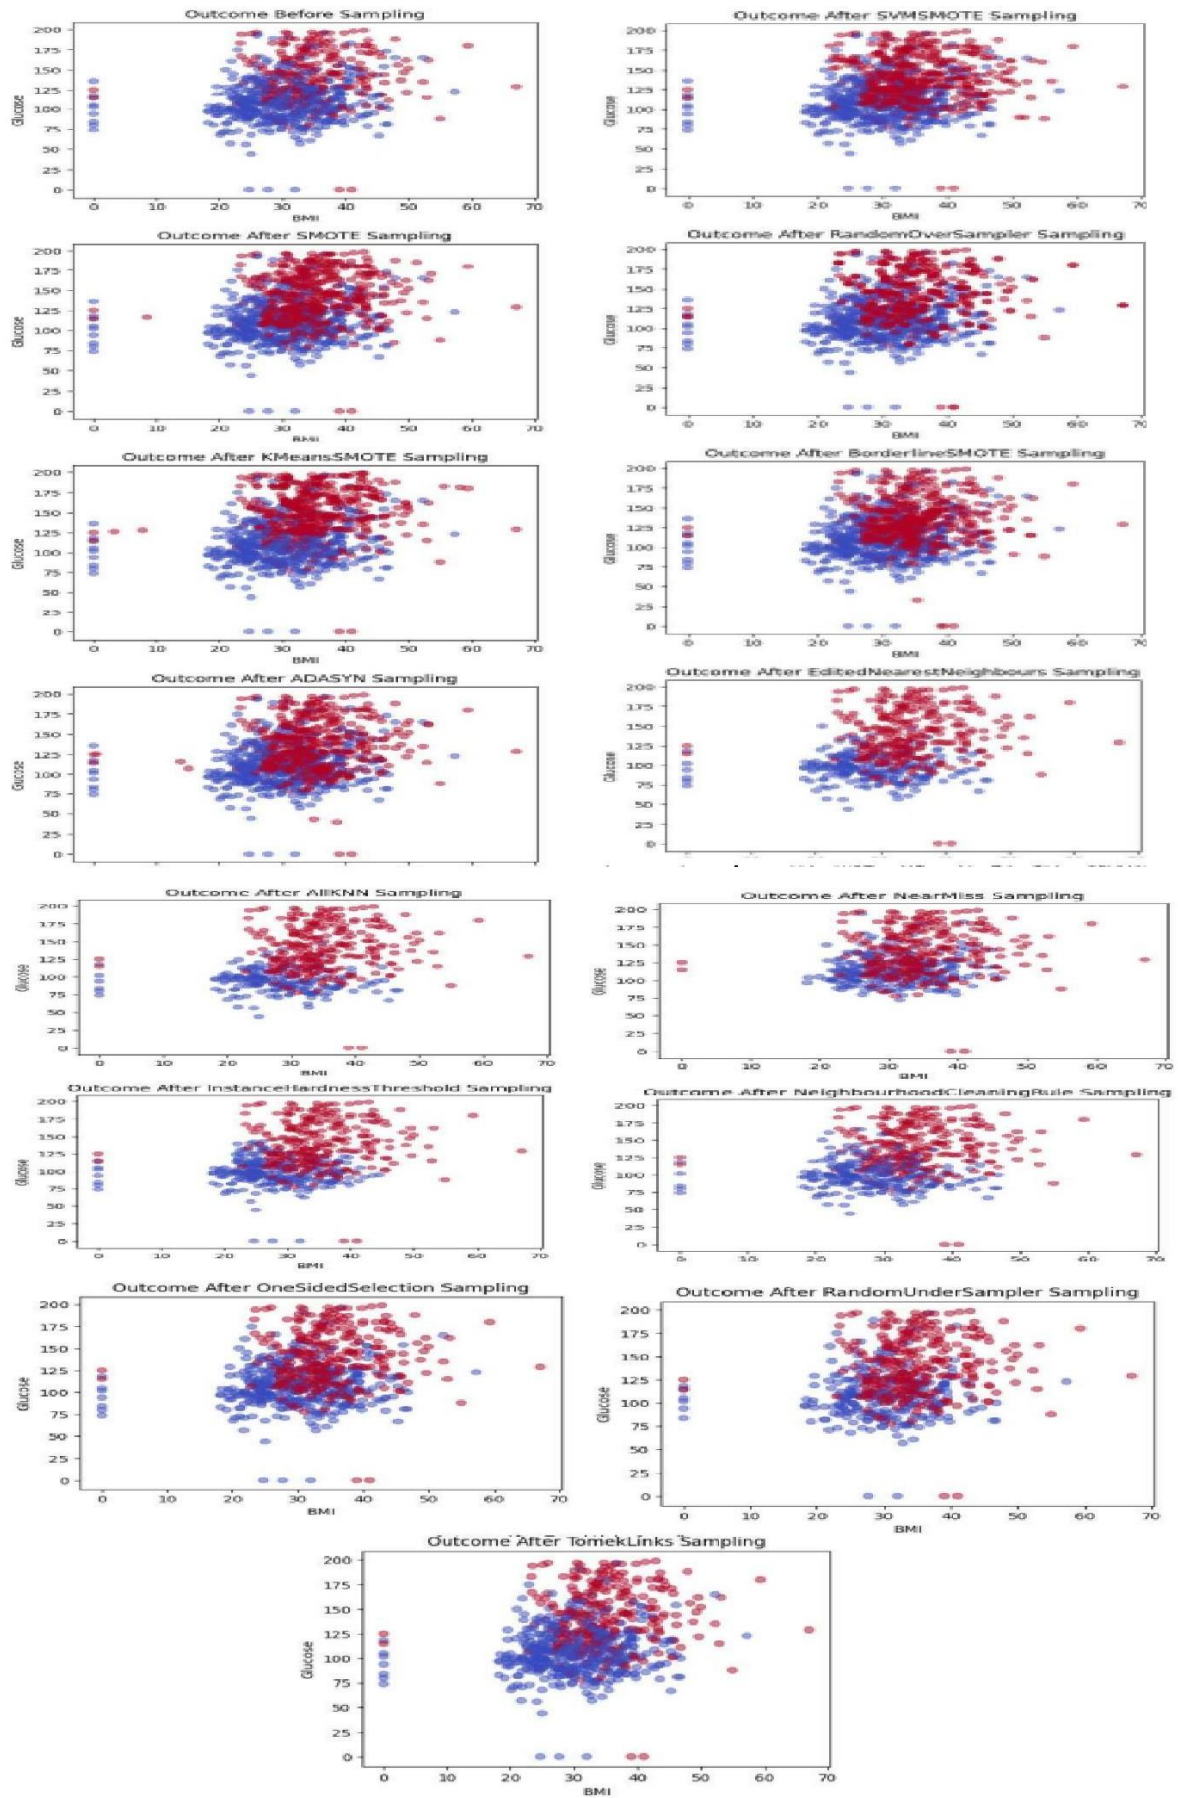


Image 2.2

When methods for correcting class imbalance are applied, the distribution of classes 0 and 1 acquired a structure similar to the example in Figure 2.1. In the subsequent stage, the new distributions of the outcome feature, as seen in Figure 2.2, are examined for each applied method, utilizing information derived from BMI and Glucose features. It was observed that the relevant methods generally did not lead to information loss.

3. Deciding whether to build a pipeline or not

We are wondering how favorable this situation would be for us when we pipeline an oversample and an undersample method. To understand this situation, we applied the random forest classification method for the balanced datasets we obtained before, and then applied the random forest classification method again by pipeline the oversample and undersample methods with the highest machine learning metric values.

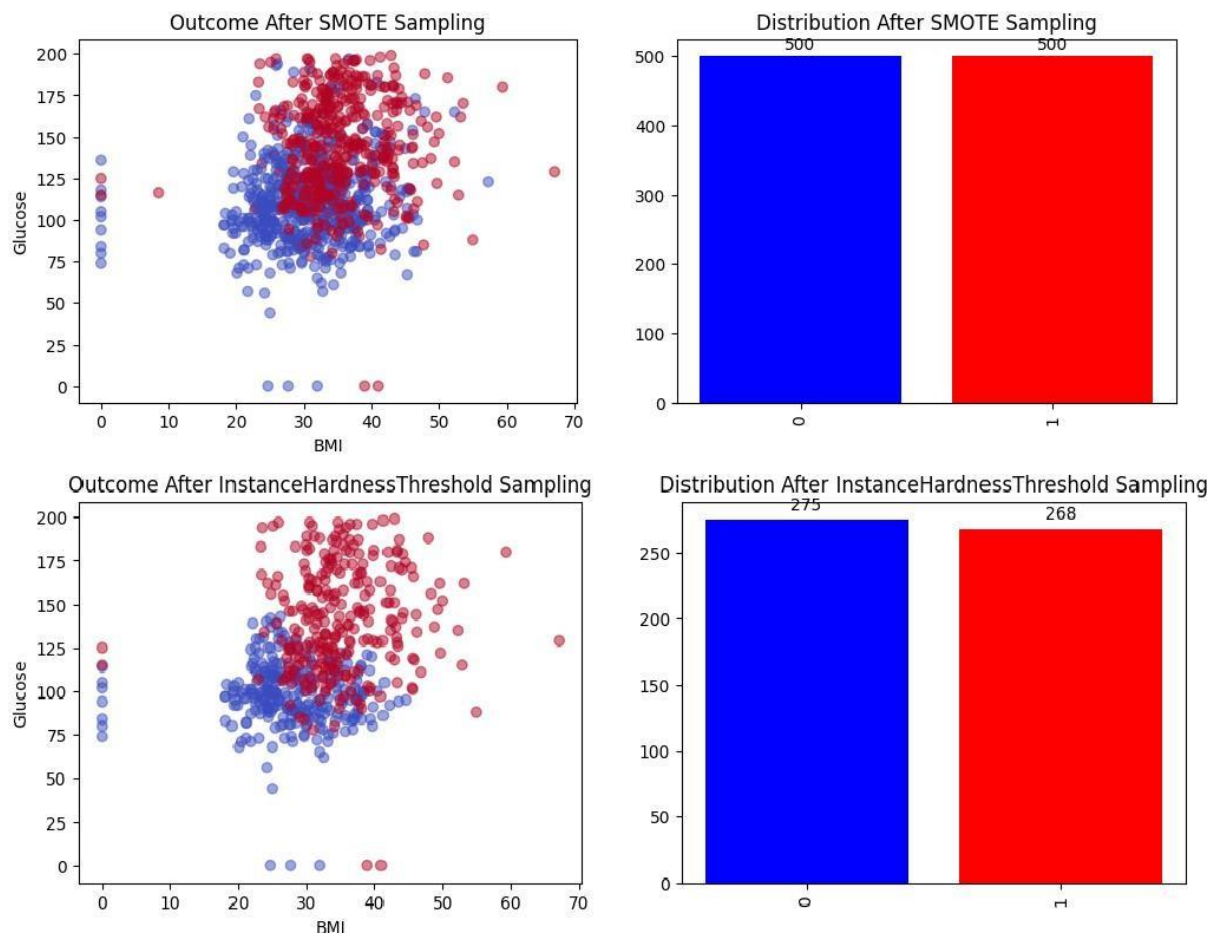


Image 3.0




Image 3.1

While the accuracy value for Smote was 0.815 and the accuracy value for Instance Hardness Threshold was 0.92660, I saw that the accuracy value decreased to 0.76 when I pipelined the two. The fact that there is no guarantee that our pipeline application will cause an increase in machine learning metrics and that I have already implemented 14 different data balance methods made us decide not to pipeline.

4. Traditional Machine Learning Algorithms

This section includes classification studies carried out on the PIMA diabetes data set with 4 different machine learning techniques (KNN, Random Forest, XGBoost, LightGBM) for the early diagnosis of diabetes. The main purpose of classification studies is to increase prediction accuracy. In this study, 14 different resampling methods are used on the data set other than the original data set to increase the success of the classifiers. For each machine learning model, 60 operations are performed without sampling and with resampling.

We used 4 different metrics to examine each classification method I applied, these are; accuracy, precision, recall and f1 score.

Accuracy: Specifies the proportion of the total number of samples that the model predicts correctly. High accuracy indicates that the overall performance of the model is good.

Precision: Specifies the probability that the samples predicted as positive by the model are actually positive. High precision indicates that the model is less likely to make false positive predictions.

Recall : Indicates how much of the truly positive samples are correctly detected by the model. A high recall indicates that the model is not missing true positives.

F1 Score: It is the harmonic average of precision and recall values. Provides a balanced measure of performance. A high F1 score indicates a model with high both precision and recall.

1. For Random Forest Classification

Sampler	Accuracy	Precision	Recall	F1 Score
Original	0.7207792207792207	0.6071428571428571	0.6181818181818182	0.6126126126126126
SMOTE	0.815	0.7711864406779662	0.900990099009901	0.8310502283105023
KMeansSMOTE	0.7960199004975125	0.8137254901960784	0.7904761904761904	0.8019323671497584
ADASYN	0.7846153846153846	0.7894736842105263	0.7731958762886598	0.78125
SVMsSMOTE	0.81	0.7889908256880734	0.8514851485148515	0.819047619047619
RandomOverSampler	0.815	0.7807017543859649	0.8811881188118812	0.8279069767441861
BorderlineSMOTE	0.765	0.75	0.801980198019802	0.7751196172248804
EditedNearestNeighbours	0.8921568627450981	0.8846153846153846	0.9019607843137255	0.8932038834951457
AllKNN	0.8829787234042553	0.8909090909090909	0.9074074074074074	0.8990825688073394
InstanceHardnessThreshold	0.926605504587156	0.9148936170212766	0.9148936170212766	0.9148936170212766
NearMiss	0.7407407407407407	0.6721311475409836	0.8367346938775511	0.7454545454545455
NeighbourhoodCleaningRule	0.8557692307692307	0.8305084745762712	0.9074074074074074	0.8672566371681415
OneSidedSelection	0.7608695652173914	0.6521739130434783	0.6382978723404256	0.6451612903225806
RandomUnderSampler	0.7870370370370371	0.75	0.7959183673469388	0.7722772272727272
TomekLinks	0.8041958041958042	0.717391304347826	0.6875	0.702127659574468

Image 4.1




Image 4.2

In general, when the Random Forest classification method is applied, the prominent dataset balancing methods are: EditedNearestNeighbours, AllKNN and InstanceHardnessThreshold. The InstanceHardnessThreshold method has higher accuracy (0.9266) and precision (0.9148) values compared to others. On the other hand, the EditedNearestNeighbours method focuses on higher recall value. The AllKNN method stands out by providing a balance between these two features.

2. For KNN Classification

Sampler	Accuracy	Precision	Recall	F1 Score
Original	0.6623376623376623	0.5245901639344263	0.5818181818181818	0.5517241379310346
SMOTE	0.74	0.7094017094017094	0.8217821782178217	0.761467889908257
KMeansSMOTE	0.7313432835820896	0.7475728155339806	0.7333333333333333	0.7403846153846153
ADASYN	0.7076923076923077	0.6851851851851852	0.7628865979381443	0.7219512195121952
SVMsSMOTE	0.77	0.7272727272727273	0.8712871287128713	0.7927927927927927
RandomOverSampler	0.73	0.7155963302752294	0.7722772277227723	0.7428571428571429
BorderlineSMOTE	0.77	0.7099236641221374	0.9207920792079208	0.8017241379310344
EditedNearestNeighbours	0.9019607843137255	0.9019607843137255	0.9019607843137255	0.9019607843137255
AllKNN	0.925531914893617	0.9607843137254902	0.9074074074074074	0.9333333333333333
InstanceHardnessThreshold	0.8807339449541285	0.8695652173913043	0.851063829787234	0.8602150537634409
NearMiss	0.6851851851851852	0.6415094339622641	0.6938775510204082	0.6666666666666666
NeighbourhoodCleaningRule	0.8461538461538461	0.8518518518518519	0.8518518518518519	0.8518518518518519
OneSidedSelection	0.7536231884057971	0.6382978723404256	0.6382978723404256	0.6382978723404256
RandomUnderSampler	0.6574074074074074	0.6111111111111112	0.673469387755102	0.6407766990291262
TomekLinks	0.7552447552447552	0.6511627906976745	0.5833333333333334	0.6153846153846155

Image 4.3




Image 4.4

In general, when the KNN classification method is applied, the prominent dataset balancing methods are: EditedNearestNeighbours, AllKNN and InstanceHardnessThreshold.

Among these, the EditedNearestNeighbours method exhibits consistent performance across accuracy (0.9019), precision (0.9019), and recall (0.9019). Notably, it emphasizes a ill-rounded approach to classification.

On the other hand, the AllKNN method demonstrates superior accuracy (0.9255) and precision (0.9608) values, showcasing its effectiveness in correctly identifying instances. It also strikes a balance with recall (0.9074), making it a robust choice for classification tasks.

In comparison, the InstanceHardnessThreshold method, while achieving slightly loir accuracy (0.8807) and precision (0.8696) values, excels in instances where recall (0.8511) is crucial.

3. For xgboost Classification

Sampler	Accuracy	Precision	Recall	F1 Score
Original	0.7077922077922078	0.5806451612903226	0.6545454545454545	0.6153846153846153
SMOTE	0.79	0.7610619469026548	0.8514851485148515	0.8037383177570093
KMeansSMOTE	0.7761194029850746	0.7884615384615384	0.780952380952381	0.7846889952153111
ADASYN	0.7435897435897436	0.7422680412371134	0.7422680412371134	0.7422680412371134
SVMsSMOTE	0.815	0.8018867924528302	0.8415841584158416	0.821256038647343
RandomOverSampler	0.805	0.7767857142857143	0.8613861386138614	0.8169014084507042
BorderlinesSMOTE	0.75	0.7256637168141593	0.8118811881188119	0.766355140186916
EditedNearestNeighbours	0.8823529411764706	0.8823529411764706	0.8823529411764706	0.8823529411764706
ALLKNN	0.9042553191489362	0.9245283018867925	0.9074074074074074	0.9158878504672898
InstanceHardnessThreshold	0.9174311926605505	0.8958333333333334	0.9148936170212766	0.9052631578947369
NearMiss	0.7129629629629629	0.6551724137931034	0.7755102040816326	0.7102803738317757
NeighbourhoodCleaningRule	0.875	0.8253968253968254	0.9629629629629629	0.8888888888888888
OneSidedSelection	0.7028985507246377	0.56	0.5957446808510638	0.577319587628866
RandomUnderSampler	0.7685185185185185	0.7068965517241379	0.8367346938775511	0.7663551401869159
TomekLinks	0.7832167832167832	0.6808510638297872	0.6666666666666666	0.6736842105263158

Image 4.5




Image 4.6

In general, when the xgboost classification method is applied, the prominent dataset balancing methods are: EditedNearestNeighbours, ALLKNN and InstanceHardnessThreshold.

EditedNearestNeighbours consistently performs well across accuracy (0.8824), precision (0.8824), and recall (0.8824), showcasing a balanced and reliable classification approach.

In contrast, ALLKNN stands out with superior accuracy (0.9043) and precision (0.9245), making it effective in correctly identifying instances. It maintains a commendable balance with recall (0.9074), presenting a robust choice for tasks prioritizing both precision and recall.

InstanceHardnessThreshold, while slightly lower in accuracy (0.9174) and precision (0.8958), excels in scenarios where recall (0.9149) is crucial. It demonstrates strength in capturing instances significant for the classification task.

In summary, EditedNearestNeighbours is consistently balanced, ALLKNN excels in accuracy and precision, and InstanceHardnessThreshold is particularly strong in recall-focused scenarios.

4. For lightgbm Classification

Sampler	Accuracy	Precision	Recall	F1 Score
Original	0.7402597402597403	0.6229508196721312	0.6909090909090909	0.6551724137931035
SMOTE	0.82	0.782608695652174	0.8910891089108911	0.8333333333333333
KMeansSMOTE	0.7910447761194029	0.8	0.8	0.8000000000000002
ADASYN	0.7333333333333333	0.7419354838709677	0.711340206185567	0.7263157894736842
SVMsSMOTE	0.79	0.7757009345794392	0.8217821782178217	0.7980769230769231
RandomOverSampler	0.825	0.7946428571428571	0.8811881188118812	0.835680751173709
BorderlineSMOTE	0.785	0.7543859649122807	0.8514851485148515	0.8
EditedNearestNeighbours	0.8725490196078431	0.8518518518518519	0.9019607843137255	0.8761904761904761
AlIcKNN	0.925531914893617	0.9607843137254902	0.9074074074074074	0.9333333333333333
InstanceHardnessThreshold	0.9174311926605505	0.88	0.9361702127659575	0.9072164948453608
NearMiss	0.7314814814814815	0.6724137931034483	0.7959183673469388	0.7289719626168225
NeighbourhoodCleaningRule	0.8653846153846154	0.8225806451612904	0.9444444444444444	0.8793103448275862
OneSidedSelection	0.717391304347826	0.5909090909090909	0.5531914893617021	0.5714285714285714
RandomUnderSampler	0.7777777777777778	0.7450980392156863	0.7755102040816326	0.76
TomekLinks	0.7972027972027972	0.7021276595744681	0.6875	0.6947368421052632

Image 4.7




Image 4.8

In general, when the lightgbm classification method is applied, the prominent dataset balancing methods are: EditedNearestNeighbours, AlIcKNN, InstanceHardnessThreshold and NeighbourhoodCleaningRule.

EditedNearestNeighbours has balanced performance across accuracy(0.87), precision(0.85), recall(0.90), and F1 score(0.87).

AlIcKNN stands out with high accuracy(0.92) and precision(0.96), effectively identifying instances while maintaining commendable balance with recall.

InstanceHardnessThreshold is slightly low in accuracy(0.91) and precision(0.88) but excels in recall-focused scenarios, showcasing strength in capturing significant instances.

NeighbourhoodCleaningRule has balanced performance with slightly low accuracy(0.86) and precision(0.82) compared to other methods, offering trade-offs.

5. Implementation of LSTM

LSTM, which stands for "Long Short-Term Memory," is a type of recurrent neural network (RNN). Recurrent neural networks are structures that can remember information from previous steps and utilize this information in future steps. However, it is known that standard RNNs struggle to handle long-term dependencies and encounter an issue called the "vanishing gradient problem." LSTM is a type of RNN designed to overcome such problems.

The ability to integrate different types of data such as BMI, glucose, age, etc., and model complex relationships has led us to implement LSTM (Long Short-Term Memory) due to its advantage in addressing long-term dependencies. LSTM allows for a more effective capture of long-term factors influencing the development and progression of the disease.

Name	Description	Value
dropout_rate	Dropout or regularization rate for LSTM layer	0.1
l1_reg	L1 regularization coefficient	0.01
epochs	Number of training epochs	30
batch_size	Batch size for training	36
test_size	Percentage of data reserved for testing	0.2

Image 5.1

Sampler	Accuracy	Precision	Recall	F1 Score
Original	0.7792207792207793	0.7560975609756098	0.5636363636363636	0.6458333333333333
SMOTE	0.745	0.7358490566037735	0.7722772277227723	0.7536231884057971
KMeansSMOTE	0.7711442786069652	0.7864077669902912	0.7714285714285715	0.7788461538461539
ADASYN	0.7230769230769231	0.7311827956989247	0.7010309278350515	0.7157894736842105
SVM SMOTE	0.73	0.7326732673267327	0.7326732673267327	0.7326732673267327
RandomOverSampler	0.705	0.7282608695652174	0.6633663366336634	0.694300518134715
BorderlineSMOTE	0.74	0.7247706422018348	0.7821782178217822	0.7523809523809524
EditedNearestNeighbours	0.8529411764705882	0.8333333333333334	0.8823529411764706	0.8571428571428571
AllKNN	0.8936170212765957	0.8928571428571429	0.9259259259259259	0.9090909090909091
InstanceHardnessThreshold	0.908256880733945	0.9302325581395349	0.851063829787234	0.8888888888888889
NearMiss	0.75	0.7291666666666666	0.7142857142857143	0.7216494845360826
NeighbourhoodCleaningRule	0.8365384615384616	0.8627450980392157	0.8148148148148148	0.838095238095238
OneSidedSelection	0.8115942028985508	0.8387096774193549	0.5531914893617021	0.6666666666666666
RandomUnderSampler	0.75	0.72	0.7346938775510204	0.7272727272727272
TomekLinks	0.8111888111888111	0.7333333333333333	0.6875	0.7096774193548386

Image 5.2




Image 5.3

A more concise explanation of the metrics obtained when examining the model's performance in a scenario where the InstanceHardnessThreshold (IHT) balancing method is used(best accuracy):

Accuracy: 0.91

It shows how much of the total samples the model correctly classified. High accuracy indicates that the overall performance of the model is good.

Precision: 0.93

It expresses the rate of examples that the model predicts as positive are actually positive. High precision indicates that positive predictions are reliable.

Recall (Sensitivity): 0.87

It measures how much of the samples it correctly classifies as truly positive. High sensitivity indicates that it does not tend to miss positive samples.

F1 Score: 0.90

As a metric that combines precision and sensitivity, it takes both false positives and false negatives into account. A high F1 score indicates a good balance achieved by the balancing method.

6. Implementation of GRU

In fact, LSTM and GRU are quite similar models. Gated Recurrent Unit (GRU) and Long Short-Term Memory (LSTM) are models of the recurrent neural network (RNN) type and are used to deal with sequential data such as time series. Both are designed to address long-term dependencies, which is the main challenge of RNN. Here are the main differences between GRU and LSTM:

1. Cell Structure:

- LSTM (Long Short-Term Memory): LSTM has a cell structure with three main gates (input, output and forgetting gates). These gates help control how much information is forgotten and how much is retained.
- GRU (Gated Recurrent Unit): GRU has a more simplified cell structure and contains only two gates: reset and update gates. Therefore, it has fewer parameters than LSTM.

2. Number of Parameters:

- LSTM: LSTM contains more parameters because there are three gates in each cell.
- GRU: GRU contains fewer parameters because the cell structure is simpler.

3. Training Duration:

- LSTM: LSTM may take longer to train because it contains more parameters.
- GRU: Since GRU contains fewer parameters, the training time may be shorter.

Parameter	Description	Value
<code>`dropout_rate`</code>	Dropout rate applied in the GRU layer.	0.1
<code>`l1_reg`</code>	L1 regularization parameter.	0.01
<code>`epochs`</code>	Number of training epochs.	30
<code>`batch_size`</code>	Batch size used during training.	36
<code>`test_size`</code>	Fraction of the data used as a test set.	0.2

Image 6.1

Sampler	Accuracy	Precision	Recall	F1 Score
Original	0.7792207792207793	0.7441860465116279	0.5818181818181818	0.6530612244897959
SMOTE	0.75	0.7428571428571429	0.7722772277227723	0.7572815533980582
KMeansSMOTE	0.7711442786069652	0.7864077669902912	0.7714285714285715	0.7788461538461539
ADASYN	0.7128205128205128	0.711340206185567	0.711340206185567	0.711340206185567
SVM SMOTE	0.74	0.7474747474747475	0.7326732673267327	0.74
RandomOverSampler	0.725	0.7395833333333334	0.7029702970297029	0.7208121827411168
BorderlineSMOTE	0.735	0.7307692307692307	0.7524752475247525	0.7414634146341463
EditedNearestNeighbours	0.8529411764705882	0.8333333333333334	0.8823529411764706	0.8571428571428571
AllKNN	0.9042553191489362	0.9090909090909091	0.9259259259259259	0.9174311926605504
InstanceHardnessThreshold	0.9174311926605505	0.9318181818181818	0.8723404255319149	0.9010989010989012
NearMiss	0.7037037037037037	0.6491228070175439	0.7551020408163265	0.6981132075471698
NeighbourhoodCleaningRule	0.8269230769230769	0.8461538461538461	0.8148148148148148	0.830188679245283
OneSidedSelection	0.8260869565217391	0.8285714285714286	0.6170212765957447	0.7073170731707318
RandomUnderSampler	0.75	0.72	0.7346938775510204	0.7272727272727272
TomekLinks	0.8111888111888111	0.7333333333333333	0.6875	0.7096774193548386

Image 6.2




Image 6.3

A more concise explanation of the metrics obtained when examining the model's performance in a scenario where the InstanceHardnessThreshold (IHT) balancing method is used(best accuracy):

Accuracy: 0.91

It shows how much of the total samples the model correctly classified. High accuracy indicates that the overall performance of the model is good.

Precision: 0.93

It expresses the rate of examples that the model predicts as positive are actually positive. High precision indicates that positive predictions are reliable.

Recall (Sensitivity): 0.87

It measures how much of the samples it correctly classifies as truly positive. High sensitivity indicates that it does not tend to miss positive samples.

F1 Score: 0.90

As a metric that combines precision and sensitivity, it takes both false positives and false negatives into account. A high F1 score indicates a good balance achieved by the balancing method.

As you can see, our result is almost the same as what I got in LSTM.

7. Applying Ensembles

Ensemble learning is a strategy of combining a number of different learning models to obtain a more powerful and generalizable model. In this section, I aimed to create an ensemble model by bringing together different neural network architectures such as CNN (Convolutional Neural Network), LSTM (Long Short-Term Memory), and GRU (Gated Recurrent Unit).

1. CNN + LSTM

Model Architecture:

- The neural network architecture consists of a combination of Conv1D (1D Convolutional) and LSTM layers.
- The data is reshaped to fit the Conv1D layer, which expects a 3D input with dimensions (samples, time steps, features).

The model includes the following layers:

- Conv1D layer with 32 filters and a kernel size of 3, using ReLU activation.
- MaxPooling1D layer with a pool size of 2.
- LSTM layer with 50 units, a kernel regularization term (L1 regularization), and return sequences set to True.
- Dropout layer with a specified dropout rate.
- Flatten layer to convert the output to a one-dimensional array.
- Dense layer with 32 units and ReLU activation.
- Output layer with 1 unit and a sigmoid activation function for binary classification.

Name	Description	Value
dropout_rate	Dropout rate for the Dropout layer in the model	0.15
l1_reg	L1 regularization parameter for the LSTM layer	0.01
epochs	Number of epochs for training the model	30
batch_size	Batch size used during training	36
test_size	Proportion of the dataset used for testing (validation)	0.2

Image 7.1

Sampler	Accuracy	Precision	Recall	F1 Score
Original	0.7532467532467533	0.6808510638297872	0.5818181818181818	0.627450980392157
SMOTE	0.725	0.7674418604651163	0.6534653465346535	0.7058823529411765
KMeansSMOTE	0.7512437810945274	0.7722772272272723	0.7428571428571429	0.7572815533980582
ADASYN	0.7538461538461538	0.7578947368421053	0.7422680412371134	0.7499999999999999
SVM SMOTE	0.745	0.7272727272727273	0.7920792079207921	0.7582938388625592
RandomOverSampler	0.74	0.7816091954022989	0.6732673267326733	0.723404255319149
BorderlineSMOTE	0.74	0.7094017094017094	0.8217821782178217	0.761467889908257
EditedNearestNeighbours	0.8921568627450981	0.8703703703703703	0.9215686274509803	0.8952380952380952
AllKNN	0.8085106382978723	0.7903225806451613	0.9074074074074074	0.8448275862068966
InstanceHardnessThreshold	0.8623853211009175	0.9210526315789473	0.7446808510638298	0.8235294117647057
NearMiss	0.6944444444444444	0.6428571428571429	0.7346938775510204	0.6857142857142857
NeighbourhoodCleaningRule	0.7884615384615384	0.7580645161290323	0.8703703703703703	0.8103448275862069
OneSidedSelection	0.7681159420289855	0.6744186046511628	0.6170212765957447	0.6444444444444444
RandomUnderSampler	0.75	0.72	0.7346938775510204	0.7272727272727272
TomekLinks	0.7622377622377622	0.6590909090909091	0.6041666666666666	0.6304347826086956

Image 7.2




Image 7.3

As can be seen, the dataset balancing method with the highest machine learning metric values in the architecture I designed is EditedNearestNeighbours.

2. CNN + GRU

Model Architecture:

- The neural network architecture includes Conv1D and GRU layers for sequence modeling.

- The data is reshaped to fit the Conv1D layer, which expects a 3D input with dimensions (samples, time steps, features).

The model includes the following layers:

- Conv1D layer with 32 filters and a kernel size of 3, using ReLU activation.
- MaxPooling1D layer with a pool size of 2.
- GRU layer with 50 units, a kernel regularization term (L1 regularization), and return sequences set to True.
- Dropout layer with a specified dropout rate.
- Flatten layer to convert the output to a one-dimensional array.
- Dense layer with 32 units and ReLU activation.
- Output layer with 1 unit and a sigmoid activation function for binary classification.

Name	Description	Value
dropout_rate	Dropout rate for the Dropout layer in the model	0.15
l1_reg	L1 regularization parameter for the GRU layer	0.01
epochs	Number of epochs for training the model	30
batch_size	Batch size used during training	36
test_size	Proportion of the dataset used for testing (validation)	0.2

Image 7.4

Sampler	Accuracy	Precision	Recall	F1 Score
Original	0.7727272727272727	0.7380952380952381	0.5636363636363636	0.6391752577319588
SMOTE	0.745	0.7717391304347826	0.7029702970297029	0.7357512953367875
KMeansSMOTE	0.7761194029850746	0.7941176470588235	0.7714285714285715	0.782608695652174
ADASYN	0.7333333333333333	0.7586206896551724	0.6804123711340206	0.717391304347826
SVM SMOTE	0.775	0.7692307692307693	0.7920792079207921	0.7804878048780488
RandomOverSampler	0.73	0.7901234567901234	0.6336633663366337	0.7032967032967031
BorderlineSMOTE	0.71	0.6902654867256637	0.7722772277227723	0.7289719626168225
EditedNearestNeighbours	0.8725490196078431	0.8392857142857143	0.9215686274509803	0.8785046728971961
AllKNN	0.8404255319148937	0.8823529411764706	0.8333333333333334	0.8571428571428571
InstanceHardnessThreshold	0.8440366972477065	0.7884615384615384	0.8723404255319149	0.8282828282828283
NearMiss	0.6944444444444444	0.625	0.8163265306122449	0.7079646017699115
NeighbourhoodCleaningRule	0.7980769230769231	0.7538461538461538	0.9074074074074074	0.8235294117647058
OneSidedSelection	0.782608695652174	0.7297297297297297	0.574468085106383	0.6428571428571429
RandomUnderSampler	0.7314814814814815	0.6666666666666666	0.8163265306122449	0.7339449541284403
TomekLinks	0.7552447552447552	0.6382978723404256	0.625	0.631578947368421

Image 7.5




Image 7.6

As can be seen, the dataset balancing method with the highest machine learning metric values in the architecture I designed is EditedNearestNeighbours.

3. CNN + LSTM + GRU

Model Architecture:

- The neural network model includes Conv1D, MaxPooling1D, LSTM, GRU, Dropout, Flatten, Dense layers.
- The input data is reshaped to fit the Conv1D layer. Conv1D layer expects 3D input: (samples, time steps, features).

The model includes the following layers:

- Conv1D layer: 32 filters, 3 kernel sizes and ReLU activation.
- MaxPooling1D layer: A pooling size of 2.
- LSTM layer: 50 units, L1 regularization term and return_sequences set to True.
- GRU layer: 50 units, L1 regularization term and return_sequences set to True.
- Dropout tier: With a certain dropout rate.
- Flatten layer: To convert the output into a one-dimensional array.
- Dense layer: 32 units and ReLU activation.
- Output layer: 1 unit and sigmoid activation function for binary classification.

Name	Description	Value
dropout_rate	Dropout rate for the Dropout layer in the model	0.15
l1_reg	L1 regularization parameter for the LSTM and GRU layers	0.01
epochs	Number of epochs for training the model	30
batch_size	Batch size used during training	36
test_size	Proportion of the dataset used for testing (validation)	0.2

Image 7.7

Sampler	Accuracy	Precision	Recall	F1 Score
Original	0.7727272727272727	0.6851851851851852	0.6727272727272727	0.6788990825688074
SMOTE	0.735	0.7352941176470589	0.7425742574257426	0.7389162561576355
KMeansSMOTE	0.7263681592039801	0.7049180327868853	0.819047619047619	0.7577092511013216
ADASYN	0.7076923076923077	0.7222222222222222	0.6701030927835051	0.6951871657754011
SVM SMOTE	0.71	0.7362637362637363	0.6633663366336634	0.6979166666666666
RandomOverSampler	0.725	0.7395833333333334	0.7029702970297029	0.7208121827411168
BorderlineSMOTE	0.735	0.7181818181818181	0.7821782178217822	0.7488151658767773
EditedNearestNeighbours	0.8921568627450981	0.8703703703703703	0.9215686274509803	0.8952380952380952
AllKNN	0.8191489361702128	0.8363636363636363	0.8518518518518519	0.8440366972477065
InstanceHardnessThreshold	0.7889908256880734	0.7068965517241379	0.8723404255319149	0.780952380952381
NearMiss	0.7222222222222222	0.6862745098039216	0.7142857142857143	0.7000000000000001
NeighbourhoodCleaningRule	0.7884615384615384	0.7580645161290323	0.8703703703703703	0.8103448275862069
OneSidedSelection	0.7753623188405797	0.6818181818181818	0.6382978723404256	0.6593406593406593
RandomUnderSampler	0.7314814814814815	0.6923076923076923	0.7346938775510204	0.7128712871287128
TomekLinks	0.7342657342657343	0.6	0.625	0.6122448979591836

Image 7.8

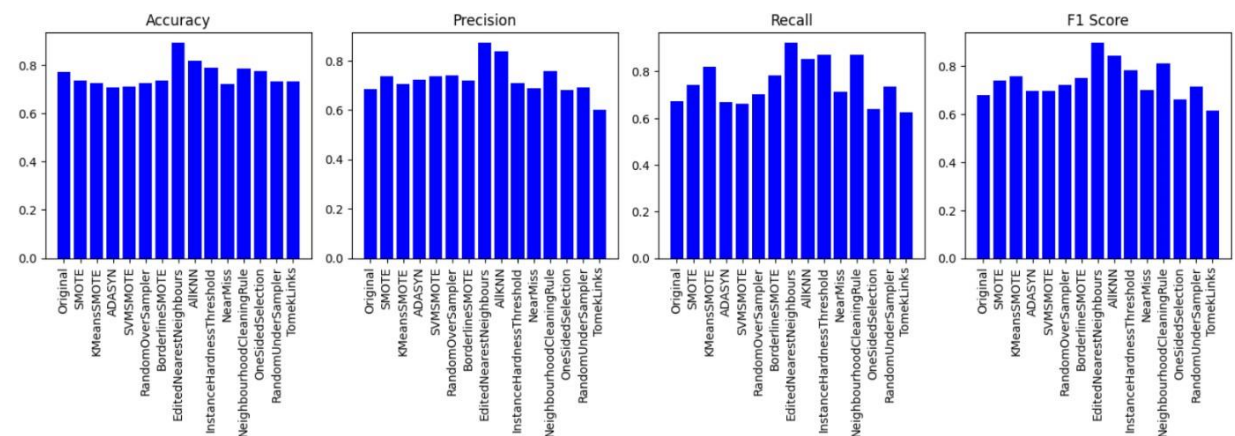


Image 7.9

As can be seen, the dataset balancing method with the highest machine learning metric values in the architecture I designed is EditedNearestNeighbours.

8. Comparing All Algorithms and Recommending Model

Algorithm	Balancing Method	Accuracy	Precision	F1 Score
Random Forest	InstanceHardnessThreshold	0.9266	0.9148	0.9148
KNN	AIKNN	0.9255	0.9607	0.9333
Xgboost	InstanceHardnessThreshold	0.9174	0.8958	0.9052
Lightgbm	AIKNN	0.9255	0.9607	0.9333
LSTM	InstanceHardnessThreshold	0.9082	0.9302	0.8888
GRU	InstanceHardnessThreshold	0.9174	0.9318	0.9010
CNN + LSTM	EditedNearestNeighbours	0.8921	0.8703	0.8952
CNN + GRU	EditedNearestNeighbours	0.8725	0.8392	0.8785
CNN + LSTM + GRU	EditedNearestNeighbours	0.8921	0.8703	0.8952

Image 8.1

Table 8.1 shows the balancing methods and metric values that give the highest machine learning algorithms in each algorithm I applied.

Machine Learning metrics are critical measurements used to evaluate the performance of the model. Confusion matrix, accuracy and loss values, which are among these metrics, are important tools to evaluate the classification ability of the model and the training process.

The confusion matrix is a table that shows how correctly or incorrectly the model classifies between classes. This matrix allows us to analyze the performance of the model in detail, especially in multi-class problems.

Accuracy refers to the ratio of correctly classified samples to the total number of samples. This metric is used as a measure of overall performance. Loss is a value that measures how far the model's predictions are from the actual values. Lower loss values indicate that the model performs better.

Graphs showing the change of these metrics over time help us understand how the model improves as it is trained. For example, the accuracy over training graph shows how the model's accuracy value changes during the training process, while the loss over training graph shows how the model's loss value decreases or increases.


The use of these metrics plays a critical role in the process of assessing the reliability of the model and optimizing its performance. To put it from an academic perspective, these metrics are basic tools for making a healthy performance analysis during the training and evaluation phases of the model.

If I give general information about the confusion matrix;

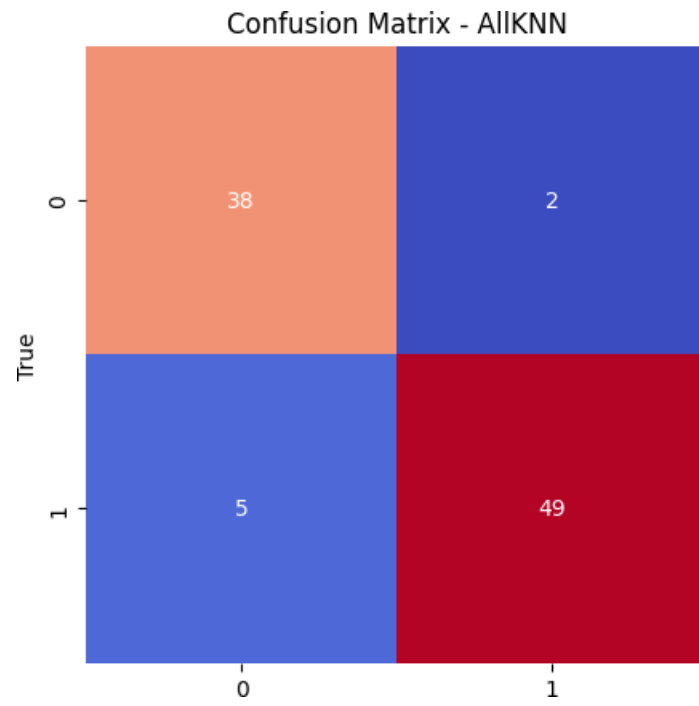
- True Positive (TP): Accurately identifying true diabetics.
- True Negative (TN): Accurately detecting those who do not actually have diabetes.
- False Positive (FP): Incorrectly identifying people who do not actually have diabetes as having diabetes.
- False Negative (FN): People who actually have diabetes are incorrectly identified as not having diabetes.

	Diyabet Değil	Diyabet
Gerçek Diyabet Değil	(True Negative)	(False Positive)
Gerçek Diyabet	(False Negative)	(True Positive)


Image 8.2




The Best in Random Forest






The Best in KNN






The Best in Xgboost






The Best in Lightgbm






The Best in LSTM




The Best in GRU



The Best in CNN+LSTM




The Best in CNN+GRU



The Best in CNN+LSTM+GRU

When both the numerical machine learning metrics and the images above are examined, no serious difference is observed between the obtained algorithms. Due to this situation, I hesitate to say that it is more advantageous to use exactly this algorithm. However, despite this situation, I can say that the metric results obtained in Random Forest are quite efficient in diagnosing the disease. For this reason, the algorithm I recommend in diagnosing diabetes is to apply Random Forest after applying the InstanceThreshold balancing method to the data in this project.



Recommended Algorithm For Detecting Diabetes

9. Results

This study aims to develop a predictive model for measuring diabetes. A significant portion of the global population suffers from diabetes, making it a crucial health concern. The goal of this research is to create an effective prediction model for diabetes. The performance of the model is evaluated on the Pima Indian Diabetes dataset.

Various balancing methods are applied to the dataset, and several algorithms are tested. In this context, the InstanceThreshold balancing method is applied to the dataset, followed by the implementation of the Random Forest Classification algorithm. The balanced dataset, obtained by applying InstanceThreshold, is split into an 80% training set and a 20% test set, resulting in a 92.66% accuracy.

When compared to state-of-the-art models, the developed model demonstrates effective performance in predicting diabetes. Future studies aim to assist healthcare professionals in the early diagnosis of diabetes and plan for the development of more advanced applications.

References

<https://tr.wikipedia.org/wiki/Diyabet>

<https://dergipark.org.tr/tr/pub/gazimmfd/issue/72928/880750>

<https://dergipark.org.tr/en/download/article-file/1371423>

<https://machinelearningmastery.com/smote-oversampling-for-imbalanced-classification/>